

# Add some SPARK to your ETL Pipeline

---

ANDREW KARCHER



Traditional vs. Modern ETL

What is Spark?

What is Azure Databricks?

Developing with Databricks

# Agenda

# What is Traditional ETL?



Vendors such as Talend, SQL Server Integration Services, Informatica, etc.



Focused on periodic batch based processes



Heavily dependent on fixed schemas



Typically a Scale-Up to handle increased loads



ETL = Extract, Transform, Load

# What has changed?



Business Intelligence -> Big Data



Relational Databases -> Files, APIs



Batch -> Streaming



Applications -> Microservices



Daily Updates -> Real Time

# What is Spark?

---

Unified  
Analytics Engine

Developed at  
AmpLab at UC  
Berkeley in 2009

100% Open  
Source

Multi Language  
Support

Optimized for  
Scale Out  
Processing

## Why Spark?

---

Need for better, faster processing than Hadoop

---

Industry shift to more text based storage

---

Need to have scale out data processing technology

---

Processing layer distinct from the Storage Layer

# Spark Ecosystem

---



# What is Azure Databricks?

---



Azure Databricks is an Apache Spark-based analytics platform optimized for the Microsoft Azure cloud services platform. Designed with the founders of Apache Spark, Databricks is integrated with Azure to provide one-click setup, streamlined workflows, and an interactive workspace that enables collaboration between data scientists, data engineers, and business analysts.



# What is Databricks?

---



Founded in 2013 by the  
original developers of Spark



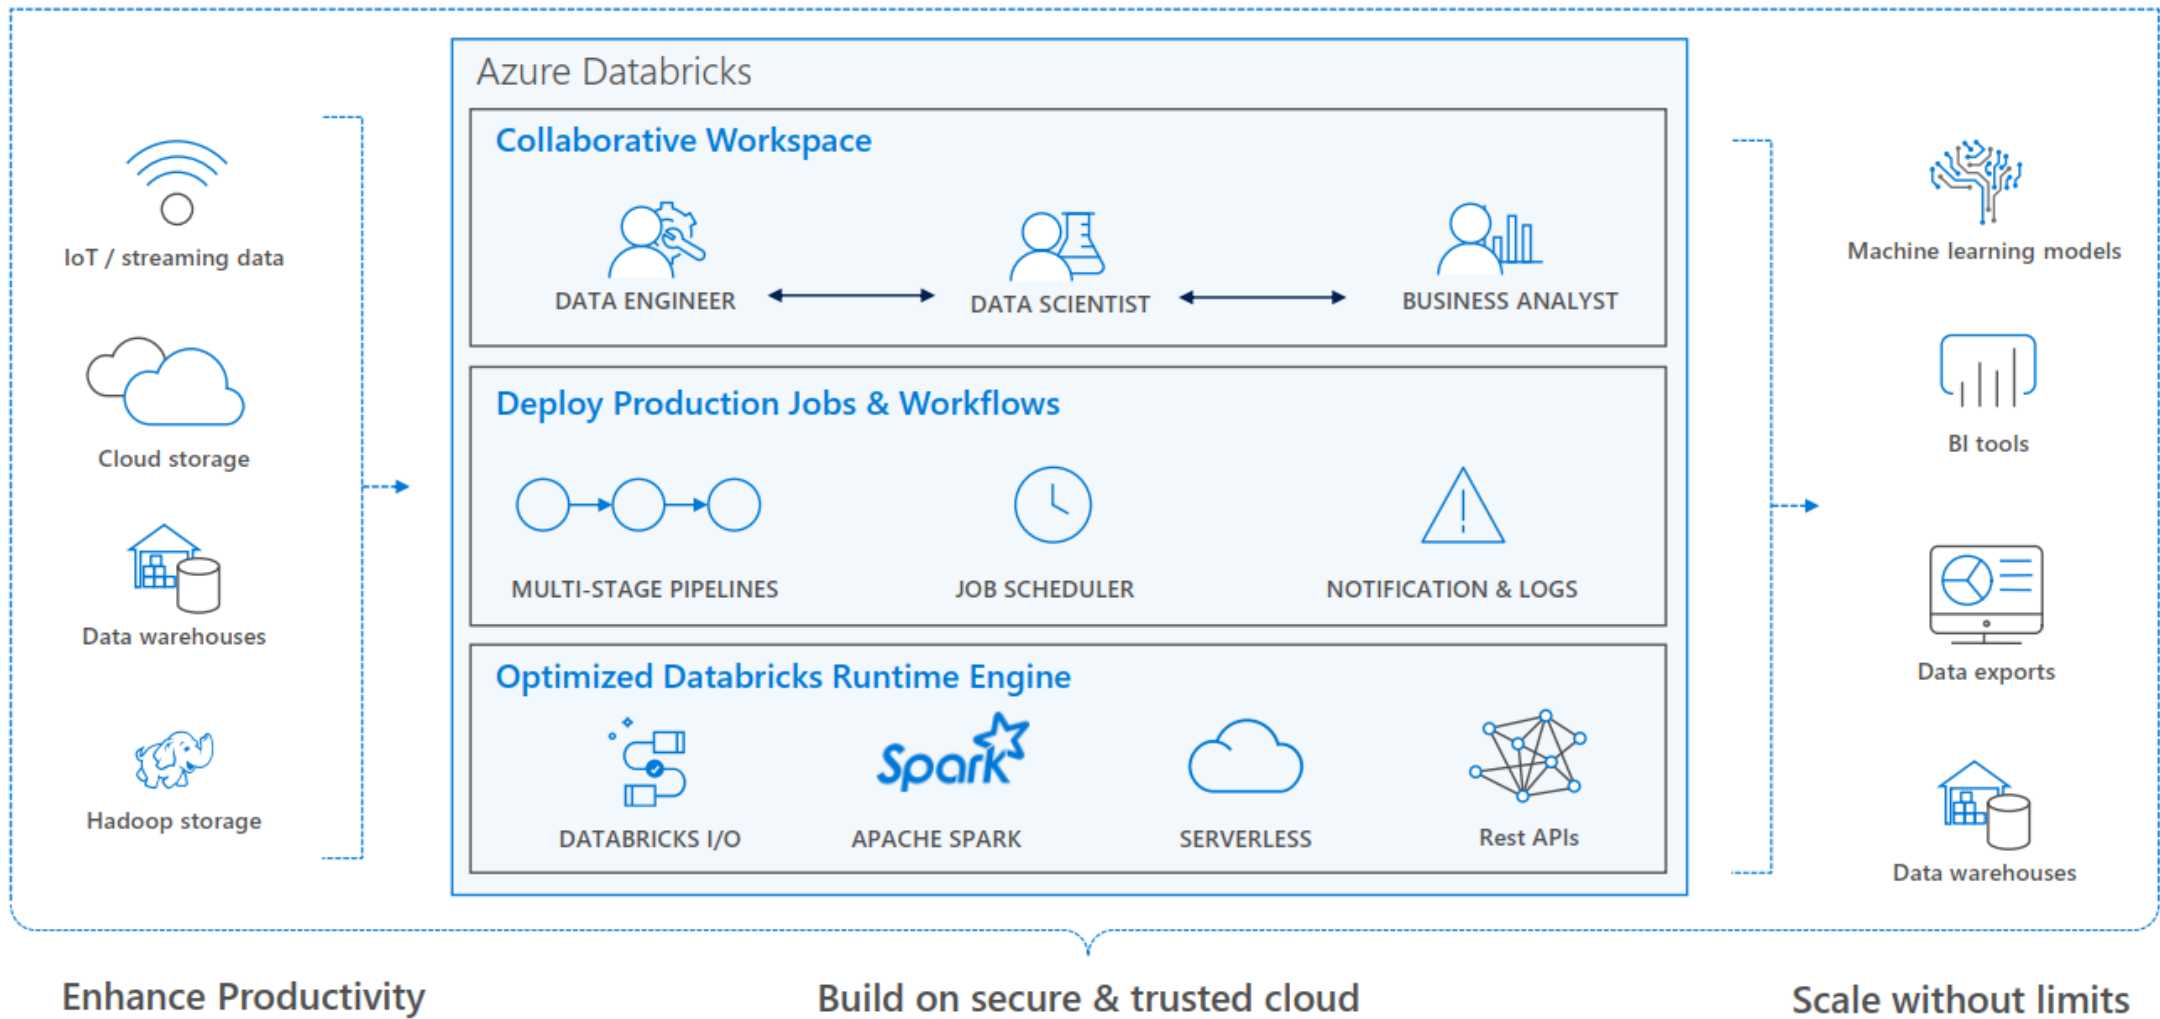
Works within the  
community to help support  
open source development



Developed Databricks Cloud  
Hosted Platform



Goal is to enable customers  
to be successful with Big  
Data



# Collaborative Workspace

Single Click to Launch Environment

Interactive Exploration with  
Notebooks

Collaborate with colleagues as well as  
integrate with Source Control

In-Notebook Visualizations

# Deployment



Jobs with Ability to Schedule



Create Multi-Stage Pipelines with Programming Control Structures



Turn Notebooks or JARs into Resilient Spark Jobs



Alerts and Audit Logs



Native Integration with other Azure Services

# Runtime



Fully Managed Cloud Platform  
(Azure, AWS)




Serverless and Elastic Cloud



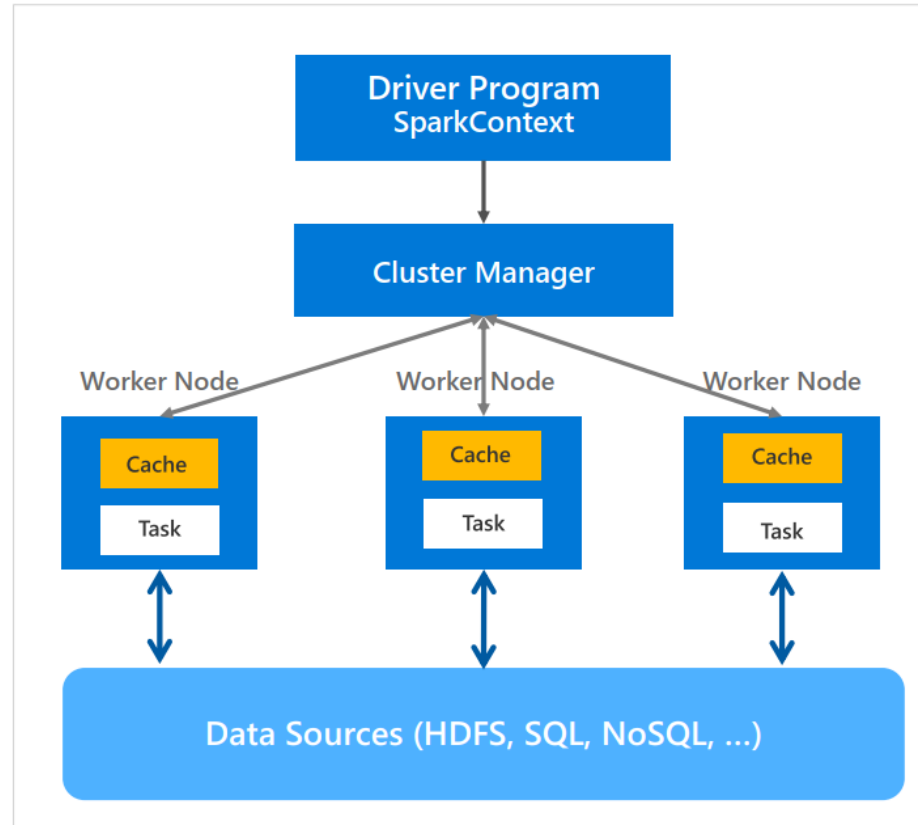
Able to operate and Scale to  
Massive Scale

# Azure Core Artifacts

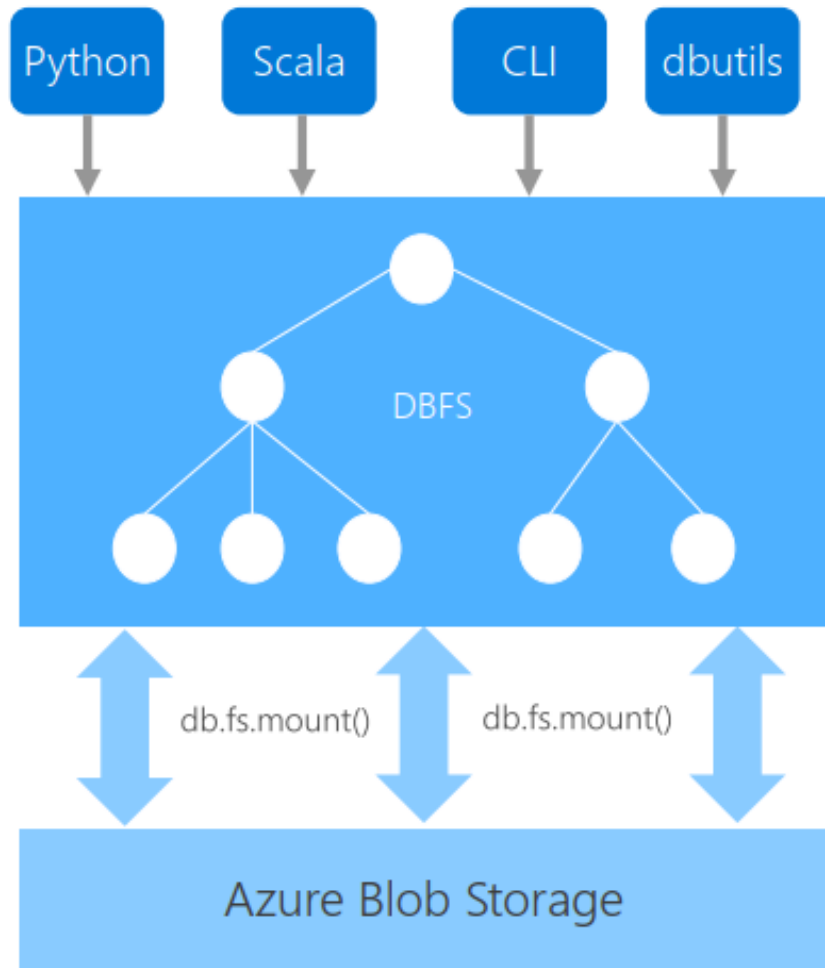
---



# General Spark Cluster Architecture







# Databricks File System

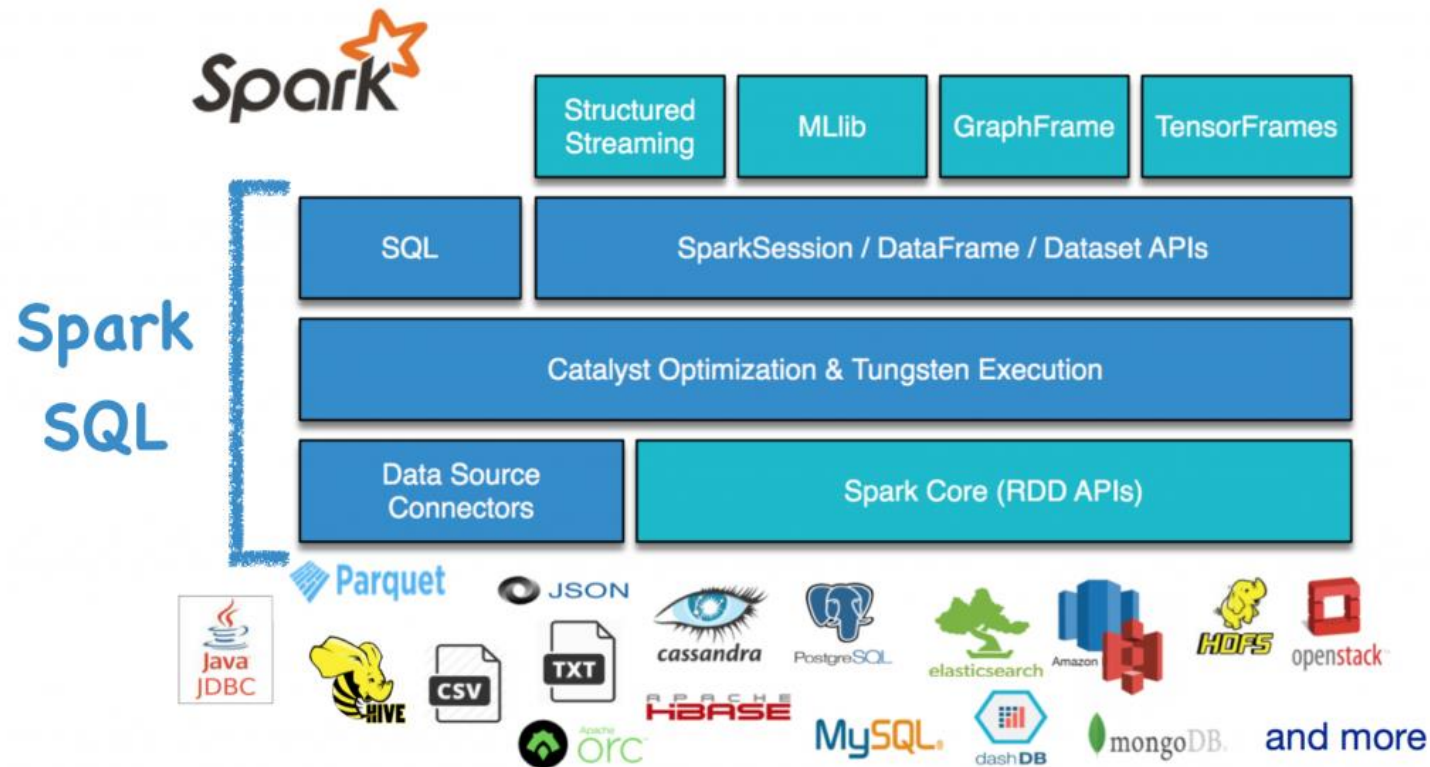
Azure Storage Buckets can be mounted in DBFS for all users to use without keys

Data is persisted in File Storage so data is not lost when cluster is removed

Pre-Installed on Databricks Clusters

Data can be cached on worker nodes on SSDs

# Spark SQL Overview



Can Query data in a wide variety of data sources in one query language

Can be queried using SQL or HiveQL

Bindings in Python, Scala, and Java

Built-In Support for Structured Streaming

# Spark Databases and Tables



Databases are a collection of related tables



Tables are defined using the GUI or programmatically using APIs or Notebooks



Utilizes the Hive Metastore to manage tables



Like a Dataframe, any Spark operation can be applied to Tables (Filtering, Caching, etc.)



Spark SQL is able to dynamically generate partitions at the file storage level

# Spark Demo

---

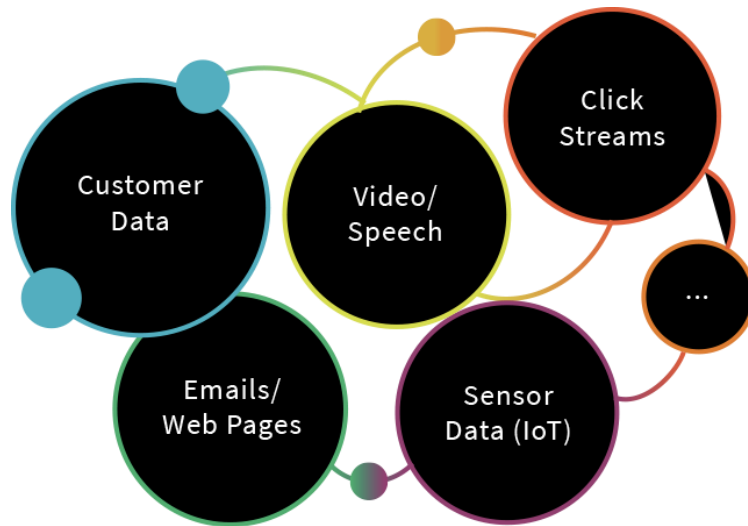
# Databricks Delta

---

AKA “DELTA LAKE”

# Data Lakes - A Key Enabler of Analytics

---



## Data Science and ML



- Recommendation Engines
- Risk, Fraud, & Intrusion Detection
- Customer Analytics
- IoT & Predictive Maintenance
- Genomics & DNA Sequencing



Failed production jobs leave data in corrupt state requiring tedious recovery



Lack of schema enforcement creates inconsistent and low quality data



Lack of consistency makes it almost impossible to mix appends and reads, batch and streaming

## Data Reliability Challenges

# Performance Challenges


---



Too many small or very big files - more time opening & closing files rather than reading contents (worse with streaming)



Partitioning aka “poor man’s indexing” - breaks down if you picked the wrong fields or when data has many dimensions, high cardinality columns



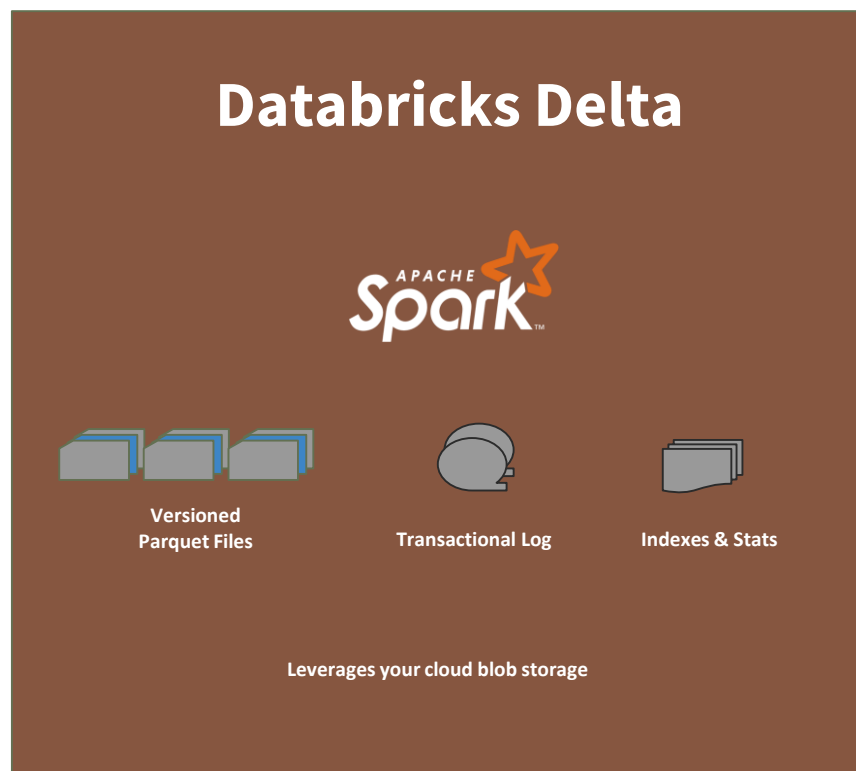
No caching - cloud storage throughput is low (S3 is 20-50MB/s/core vs 300MB/s/core for local SSDs)



# Databricks Delta

Next-generation engine built on top of Spark

---



Co-designed compute & storage

Compatible with Spark API's


Built on open standards  
(Parquet)

# Features of Databricks Delta

---



ACID  
TRANSACTIONS



SCHEMA  
ENFORCEMENT



UPSERTS



DATA  
VERSIONING

# Databricks Delta Performance

Compaction

Caching

Data Skipping

Z-Order Indexes

Optimized Parquet

# Demo: Delta and Streaming

---

## Conclusion: Should you add Spark to your ETL Pipeline? (Pros)

As you move from a Batch based pipeline to Streaming it is worth a consideration

If you want to tap into ML Pipelines it is definitely something to consider

I don't think you need to migrate all your existing ETL pipelines to Spark

If you are doing your data work in the cloud then it is definitely something to consider

Conclusion:  
Should you add  
Spark to your  
ETL Pipelines?  
(Cons)

Do you have the skills in-house? If not, are you willing to invest in training your existing people or hiring new ones?

Is there a business reason that supports investing in a new technology?

Is your data in the cloud or on-premise?  
(Can you manage a Spark Cluster internally?)

# About Me

<http://www.andrewkarcher.com>

Twitter: @akarcher

LinkedIn, Twitter

Email: [akarcher@gmail.com](mailto:akarcher@gmail.com)

Company – Data Architect, Becton Dickinson (BD)

## Community Work

- Former Organizer for SQL Saturday OC
- President San Diego .Net User Group
- Organizer – Data Engineering San Diego
- Former SQL Server MVP

# Questions??

[AKARCHER@GMAIL.COM](mailto:AKARCHER@GMAIL.COM)

@AKARCHER

[HTTP://WWW.ANDREWKARCHER.COM](http://WWW.ANDREWKARCHER.COM)